

```
In [1]: 1 #!pip install geopy
```

```
In [2]: 1 #!pip install geocoder
```

```
In [3]: 1 #!pip install uszipcode
```

```
In [4]: 1 #!pip install arcgis
```

```
In [5]: 1 #!pip install --force-reinstall numpy==1.23.3
```

```
In [6]: 1 #import numpy  
2 #numpy.version.version
```

```
In [7]: 1 #!pip install numba
```

```
In [856]: 1 #!pip install gmaps  
2
```

...

```
In [790]: 1 import pandas as pd  
2 import numpy as np  
3 import seaborn as sns  
4 import matplotlib.pyplot as plt  
5 import seaborn as sns  
6 from sklearn.model_selection import train_test_split, cross_validate, Shuffle  
7  
8 from sklearn.metrics import ConfusionMatrixDisplay, confusion_matrix, roc_auc  
9 plot_confusion_matrix, precision_recall_curve  
10  
11 from sklearn.preprocessing import OneHotEncoder, StandardScaler  
12 from sklearn.impute import SimpleImputer  
13 from sklearn.pipeline import Pipeline  
14 from sklearn.compose import ColumnTransformer  
15  
16 from sklearn.linear_model import LogisticRegression  
17 from sklearn.neighbors import KNeighborsClassifier  
18 from sklearn.tree import DecisionTreeClassifier  
19  
20 from sklearn.ensemble import RandomForestClassifier, VotingClassifier, AdaBo  
21 AdaBoostClassifier, GradientBoostingClassifier  
22  
23 from sklearn.cluster import KMeans  
24 import xgboost  
25 #import geopy  
26 #from geopy.geocoders import Nominatim  
27 #from arcgis.geocoding import reverse_geocode  
28
```

Making GPU as processor

```
In [9]: 1 import numba
        2 numba.__version__
```

Out[9]: '0.56.2'

```
In [10]: 1 from numba import jit, cuda
        2 import numpy as np
        3 # to measure exec time
        4 from timeit import default_timer as timer
        5
        6 # normal function to run on cpu
        7 def func(a):
        8     for i in range(10000000):
        9         a[i] += 1
       10
       11 # function optimized to run on gpu
       12 #@jit(target_backend='cuda')
       13 def func2(a):
       14     for i in range(10000000):
       15         a[i] += 1
       16 if __name__ == "__main__":
       17     n = 10000000
       18     a = np.ones(n, dtype = np.float64)
       19
       20     start = timer()
       21     func(a)
       22     print("without GPU:", timer()-start)
       23
       24     start = timer()
       25     func2(a)
       26     print("with GPU:", timer()-start)
```

without GPU: 2.4438747999999997

with GPU: 2.4060063999999999


Importing & Exploring Data

```
In [11]: 1 df_crashes = pd.read_csv('data/Traffic_Crashes_-_Crashes.csv')
         2 df_crashes.head()
```

```
Out[11]:
```

	CRASH_RECORD_ID	RD_NO	CRASH_DATE_EST_I	CRASH_DA
0	062f5a6f6b87b762165d4da04d6d3a181385776a10b051...	JF378246	NaN	08/31/20 10:13:00 F
1	0115ade9a755e835255508463f7e9c4a9a0b47e9304238...	JF318029	NaN	07/15/20 12:45:00 A
2	017040c61958d2fa977c956b2bd2d6759ef7754496dc96...	JF324552	NaN	07/15/20 06:50:00 F
3	01aaa759c6bbefd0f584226fbd88bdc549de3ed1e46255...	JF319819	NaN	07/15/20 05:10:00 F
4	04f21d51f8189e34abf37c7973607fa076965d216b514f...	JC366684	NaN	07/22/20 12:00:00 F

5 rows × 49 columns



```
In [12]: 1 df_crashes.info()
```

...

```
In [13]: 1 df_crashes.TRAFFIC_CONTROL_DEVICE.value_counts()
```

```
Out[13]: NO CONTROLS          376542
TRAFFIC SIGNAL          181633
STOP SIGN/FLASHER       64989
UNKNOWN                 22416
OTHER                   4205
LANE USE MARKING        1226
YIELD                   913
OTHER REG. SIGN         683
OTHER WARNING SIGN      565
RAILROAD CROSSING GATE  423
PEDESTRIAN CROSSING SIGN 357
DELINEATORS            247
POLICE/FLAGMAN          238
SCHOOL ZONE             236
FLASHING CONTROL SIGNAL 223
OTHER RAILROAD CROSSING 153
RR CROSSING SIGN        78
NO PASSING              36
BICYCLE CROSSING SIGN   19
Name: TRAFFIC_CONTROL_DEVICE, dtype: int64
```

```
In [14]: 1 df_crashes.isna().sum()
```

```
Out[14]: CRASH_RECORD_ID          0
RD_NO          4587
CRASH_DATE_EST_I  605487
CRASH_DATE          0
POSTED_SPEED_LIMIT  0
TRAFFIC_CONTROL_DEVICE  0
DEVICE_CONDITION    0
WEATHER_CONDITION    0
LIGHTING_CONDITION    0
FIRST_CRASH_TYPE    0
TRAFFICWAY_TYPE      0
LANE_CNT          456191
ALIGNMENT           0
ROADWAY_SURFACE_COND  0
ROAD_DEFECT         0
REPORT_TYPE        17491
CRASH_TYPE          0
INTERSECTION_RELATED_I  505314
NOT_RIGHT_OF_WAY_I    624336
HIT_AND_RUN_I        453020
DAMAGE              0
DATE_POLICE_NOTIFIED  0
PRIM_CONTRIBUTORY_CAUSE  0
SEC_CONTRIBUTORY_CAUSE  0
STREET_NO           0
STREET_DIRECTION     4
STREET_NAME          1
BEAT_OF_OCCURRENCE    5
PHOTOS_TAKEN_I       647098
STATEMENTS_TAKEN_I    641669
DOORING_I            653127
WORK_ZONE_I          651287
WORK_ZONE_TYPE        652126
WORKERS_PRESENT_I     654185
NUM_UNITS            0
MOST_SEVERE_INJURY    1401
INJURIES_TOTAL        1390
INJURIES_FATAL        1390
INJURIES_INCAPACITATING  1390
INJURIES_NON_INCAPACITATING  1390
INJURIES_REPORTED_NOT_EVIDENT  1390
INJURIES_NO_INDICATION  1390
INJURIES_UNKNOWN      1390
CRASH_HOUR            0
CRASH_DAY_OF_WEEK     0
CRASH_MONTH           0
LATITUDE             4074
LONGITUDE             4074
LOCATION               4074
dtype: int64
```

```
In [15]: 1 # Keeping relevant Features
2 df_crashes_drop = df_crashes [[
3     'CRASH_RECORD_ID',
4     # 'RD_NO',
5     'CRASH_DATE',
6     'POSTED_SPEED_LIMIT',
7     'WEATHER_CONDITION',
8     # 'LIGHTING_CONDITION',
9     # 'FIRST_CRASH_TYPE',
10    'ROADWAY_SURFACE_COND',
11    'ROAD_DEFECT',
12    # 'CRASH_TYPE',
13    # 'DAMAGE',
14    # 'PRIM_CONTRIBUTORY_CAUSE',
15    # 'STREET_NAME',
16    # 'NUM_UNITS',
17    'INJURIES_TOTAL',
18    'INJURIES_FATAL',
19    'CRASH_HOUR',
20    'CRASH_DAY_OF_WEEK',
21    'CRASH_MONTH',
22    'LATITUDE',
23    'LONGITUDE',
24    # 'LOCATION',
25 ]]
26 print(df_crashes_drop.shape)
```

(655182, 13)

```
In [16]: 1 df_crashes_drop.describe()
```

```
Out[16]:
```

	POSTED_SPEED_LIMIT	INJURIES_TOTAL	INJURIES_FATAL	CRASH_HOUR	CRASH_DAY_OF_WEEK
count	655182.000000	653792.000000	653792.000000	655182.000000	655182.000000
mean	28.356959	0.184946	0.001169	13.222685	4.000000
std	6.296888	0.558173	0.037166	5.549903	1.000000
min	0.000000	0.000000	0.000000	0.000000	1.000000
25%	30.000000	0.000000	0.000000	9.000000	2.000000
50%	30.000000	0.000000	0.000000	14.000000	4.000000
75%	30.000000	0.000000	0.000000	17.000000	6.000000
max	99.000000	21.000000	4.000000	23.000000	7.000000

```
In [17]: 1 # Dropping Rows with Latitude & Longitude = 0
2 df_crashes_drop = df_crashes_drop[df_crashes_drop['LATITUDE'] != 0]
```

```
In [18]: 1 df_crashes_drop['has_injuries'] = df_crashes_drop.INJURIES_TOTAL.apply(lambda x: x > 0)
2 df_crashes_drop['has_fatality'] = df_crashes_drop.INJURIES_FATAL.apply(lambda x: x > 0)
```

```
In [1025]: 1 # Visualizing the distribution of accidents by having injuries ( Alternative
2 crash_df_ = df_crashes_drop.groupby(by=['LONGITUDE', 'LATITUDE']).agg(crashes
3 crash_df_
```

```
Out[1025]:
```

	LONGITUDE	LATITUDE	crashes	has_injuries
0	-87.936193	41.960822	1	0
1	-87.935877	41.960761	1	0
2	-87.934763	41.960230	3	0
3	-87.934510	42.008051	1	0
4	-87.934014	41.959123	1	0
...
260229	-87.524674	41.702590	8	1
260230	-87.524646	41.698928	1	0
260231	-87.524640	41.703323	1	0
260232	-87.524589	41.702571	4	1
260233	-87.524587	41.703272	7	1

260234 rows × 4 columns

```
In [862]: 1 #!pip install plotly==5.8.0
```

Collecting plotly==5.8.0
 Downloading plotly-5.8.0-py2.py3-none-any.whl (15.2 MB)
 Collecting tenacity>=6.2.0
 Downloading tenacity-8.1.0-py3-none-any.whl (23 kB)
 Installing collected packages: tenacity, plotly
 Attempting uninstall: plotly
 Found existing installation: plotly 4.11.0
 Uninstalling plotly-4.11.0:
 Successfully uninstalled plotly-4.11.0
 Successfully installed plotly-5.8.0 tenacity-8.1.0

```
In [866]: 1 import folium
2 from folium.plugins import HeatMap
```

```
In [1028]: 1 lats longs_weight = list(map(list, zip(crash_df_["LATITUDE"],
2 crash_df_["LONGITUDE"],
3 crash_df_["crashes"]
4 )
5 )
6 )
```

```
In [1029]: 1 map_obj = folium.Map(location = [41.874389144012255, -87.668751362594],tiles
2
3 HeatMap(lats longs_weight, min_opacity=0.5,max_zoom = 40,radius=9,control=Tr
4 #folium.LayerControl().add_to(map_obj)
5 map_obj
```

Out[1029]: Make this Notebook Trusted to load map: File -> Trust Notebook

```
In [21]: 1 # Dropping unknown values
2 df_crashes_drop.replace({'UNKNOWN':np.nan}, inplace=True)
```

```
In [22]: 1 # Dropping rows with Lats and Long =0
2 df_crashes_ = df_crashes_drop.dropna(subset = 'LATITUDE', axis = 0);
```


```
In [23]: 1 df_vehicles = pd.read_csv('data/Traffic_Crashes_-_Vehicles.csv')
2 df_vehicles.head()
```

```
<ipython-input-23-fce9fee1d033>:1: DtypeWarning: Columns (19,21,40,41,42,44,48,
49,50,53,55,58,59,61,71) have mixed types. Specify dtype option on import or se
t low_memory=False.
df_vehicles = pd.read_csv('data/Traffic_Crashes_-_Vehicles.csv')
```

```
Out[23]:
```

	CRASH_UNIT_ID	CRASH_RECORD_ID	RD_NO	CRASH_DATE	L
0	829999	24ddf9fd8542199d832e1c223cc474e5601b356f1d77a6...	JD124535	01/22/2020 06:25:00 AM	
1	749947	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	JC451435	09/28/2019 03:30:00 AM	
2	749949	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	JC451435	09/28/2019 03:30:00 AM	
3	749950	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	JC451435	09/28/2019 03:30:00 AM	
4	871921	af84fb5c8d996fcd3aefd36593c3a02e6e7509eeb27568...	JD208731	04/13/2020 10:50:00 PM	

5 rows × 72 columns



In [24]: 1 df_vehicles.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1335230 entries, 0 to 1335229
Data columns (total 72 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CRASH_UNIT_ID                        1335230 non-null  int64
1   CRASH_RECORD_ID                     1335230 non-null  object
2   RD_NO                               1325924 non-null  object
3   CRASH_DATE                           1335230 non-null  object
4   UNIT_NO                             1335230 non-null  int64
5   UNIT_TYPE                           1333361 non-null  object
6   NUM_PASSENGERS                      198353 non-null   float64
7   VEHICLE_ID                          1304967 non-null  float64
8   CMRC_VEH_I                          24927 non-null    object
9   MAKE                                1304962 non-null  object
10  MODEL                               1304818 non-null  object
11  LIC_PLATE_STATE                     1187047 non-null  object
12  VEHICLE_YEAR                        1092432 non-null  float64
13  VEHICLE_DEFECT                      1304967 non-null  object
14  VEHICLE_TYPE                        1304967 non-null  object
15  VEHICLE_USE                          1304967 non-null  object
16  TRAVEL_DIRECTION                   1304967 non-null  object
17  MANEUVER                           1304967 non-null  object
18  TOWED_I                             162050 non-null   object
19  FIRE_I                              1062 non-null     object
20  OCCUPANT_CNT                       1304967 non-null  float64
21  EXCEED_SPEED_LIMIT_I               2397 non-null     object
22  TOWED_BY                           120536 non-null   object
23  TOWED_TO                           74156 non-null    object
24  AREA_00_I                           47900 non-null    object
25  AREA_01_I                           351739 non-null   object
26  AREA_02_I                           216469 non-null   object
27  AREA_03_I                           126691 non-null   object
28  AREA_04_I                           128805 non-null   object
29  AREA_05_I                           199196 non-null   object
30  AREA_06_I                           206230 non-null   object
31  AREA_07_I                           185253 non-null   object
32  AREA_08_I                           202505 non-null   object
33  AREA_09_I                           76788 non-null    object
34  AREA_10_I                           110769 non-null   object
35  AREA_11_I                           217168 non-null   object
36  AREA_12_I                           213818 non-null   object
37  AREA_99_I                           146804 non-null   object
38  FIRST_CONTACT_POINT                 1293476 non-null  object
39  CMV_ID                              14038 non-null    float64
40  USDOT_NO                           8044 non-null     object
41  CCMC_NO                             1755 non-null     object
42  ILCC_NO                             1230 non-null     object
43  COMMERCIAL_SRC                      9556 non-null     object
44  GVWR                                7988 non-null     object
45  CARRIER_NAME                       13428 non-null    object
46  CARRIER_STATE                      12645 non-null    object
47  CARRIER_CITY                       12425 non-null    object
48  HAZMAT_PLACARDS_I                   276 non-null      object
49  HAZMAT_NAME                         51 non-null       object
```

50	UN_NO	489 non-null	object
51	HAZMAT_PRESENT_I	10282 non-null	object
52	HAZMAT_REPORT_I	9972 non-null	object
53	HAZMAT_REPORT_NO	1 non-null	object
54	MCS_REPORT_I	10033 non-null	object
55	MCS_REPORT_NO	6 non-null	object
56	HAZMAT_VIO_CAUSE_CRASH_I	10130 non-null	object
57	MCS_VIO_CAUSE_CRASH_I	9946 non-null	object
58	IDOT_PERMIT_NO	793 non-null	object
59	WIDE_LOAD_I	123 non-null	object
60	TRAILER1_WIDTH	2584 non-null	object
61	TRAILER2_WIDTH	303 non-null	object
62	TRAILER1_LENGTH	2124 non-null	float64
63	TRAILER2_LENGTH	60 non-null	float64
64	TOTAL_VEHICLE_LENGTH	2577 non-null	float64
65	AXLE_CNT	3786 non-null	float64
66	VEHICLE_CONFIG	11649 non-null	object
67	CARGO_BODY_TYPE	11148 non-null	object
68	LOAD_TYPE	10655 non-null	object
69	HAZMAT_OUT_OF_SERVICE_I	9674 non-null	object
70	MCS_OUT_OF_SERVICE_I	9921 non-null	object
71	HAZMAT_CLASS	939 non-null	object

dtypes: float64(9), int64(2), object(61)
memory usage: 733.5+ MB

In [25]:

```

1  # Keeping Relevant Features
2  df_vehicles_drop = df_vehicles [[
3      'CRASH_RECORD_ID',
4      # 'RD_NO',
5      # 'CRASH_DATE',
6      'VEHICLE_ID',
7      # 'MAKE',
8      # 'MODEL',
9      # 'LIC_PLATE_STATE',
10     'VEHICLE_YEAR',
11     # 'VEHICLE_DEFECT',
12     # 'VEHICLE_TYPE',
13     # 'VEHICLE_USE',
14     # 'TRAVEL_DIRECTION',
15     'OCCUPANT_CNT',
16     # 'VEHICLE_CONFIG',
17 ]]
18 print(df_vehicles_drop.shape)

```

(1335230, 4)

```
In [26]: 1 df_vehicles.VEHICLE_CONFIG.value_counts()
```

```
Out[26]: TRACTOR/SEMI-TRAILER          4572
SINGLE UNIT TRUCK, 2 AXLES, 6 TIRES    2337
BUS                                   1805
TRUCK/TRACTOR                        944
TRUCK/TRAILER                        770
UNKNOWN HEAVY TRUCK                  673
SINGLE UNIT TRUCK, 3 OR MORE AXLES    516
TRACTOR/DOUBLES                      32
Name: VEHICLE_CONFIG, dtype: int64
```

```
In [27]: 1 df_vehicles_drop.describe()
```

```
Out[27]:
```

	VEHICLE_ID	VEHICLE_YEAR	OCCUPANT_CNT
count	1.304967e+06	1.092432e+06	1.304967e+06
mean	6.789888e+05	2.013626e+03	1.079589e+00
std	3.917114e+05	1.439261e+02	7.843770e-01
min	2.000000e+00	1.900000e+03	0.000000e+00
25%	3.400725e+05	2.006000e+03	1.000000e+00
50%	6.794450e+05	2.012000e+03	1.000000e+00
75%	1.017828e+06	2.016000e+03	1.000000e+00
max	1.358762e+06	9.999000e+03	9.900000e+01

```
In [28]: 1 # Dropping rows with vehicle year bigger than 2022 and below 1970
2 df_vehicles_drop1 = df_vehicles_drop[(df_vehicles_drop['VEHICLE_YEAR'] >= 19
```

```
In [31]: 1 # Dropping Vehicles with 0 Occupant (Parked Cars)
2 df_vehicles_w_occ = df_vehicles_drop1[df_vehicles_drop.OCCUPANT_CNT != 0.0]
3 print(df_vehicles_w_occ.shape)
```

```
(1160416, 4)
```

```
<ipython-input-31-a0910bcd00f>:2: UserWarning: Boolean Series key will be rein
dexed to match DataFrame index.
```

```
df_vehicles_w_occ = df_vehicles_drop1[df_vehicles_drop.OCCUPANT_CNT != 0.0]
```

```
In [32]: 1 df_vehicles_w_occ.isna().sum()
```

```
Out[32]: CRASH_RECORD_ID          0
VEHICLE_ID          30263
VEHICLE_YEAR        231819
OCCUPANT_CNT        30263
dtype: int64
```

In [33]:

```
1 # Replacing all UNKNOWN values to missing values
2 df_vehicles_w_occ.replace({'UNKNOWN':np.nan}, inplace=True)
```

<ipython-input-33-fdbee41cc7b6>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_vehicles_w_occ.replace({'UNKNOWN':np.nan}, inplace=True)
```

In [34]:

```
1 df_people = pd.read_csv('data/Traffic_Crashes_-_PEOPLE.csv')
2 df_people.head()
3
```


<ipython-input-34-78bd59adc723>:1: DtypeWarning: Columns (20,24,25,26,29) have mixed types. Specify dtype option on import or set low_memory=False.

```
df_people = pd.read_csv('data/Traffic_Crashes_-_PEOPLE.csv')
```

Out[34]:

	PERSON_ID	PERSON_TYPE		CRASH_RECORD_ID	RD_NO	VEH
0	O749947	DRIVER	81dc0de2ed92aa62baccab641fa377be7feb1cc47e6554...	JC451435		{
1	O871921	DRIVER	af84fb5c8d996fcd3aefd36593c3a02e6e7509eeb27568...	JD208731		{
2	O10018	DRIVER	71162af7bf22799b776547132ebf134b5b438dcf3dac6b...	HY484534		
3	O10038	DRIVER	c21c476e2ccc41af550b5d858d22aaac4ffc88745a1700...	HY484750		
4	O10039	DRIVER	eb390a4c8e114c69488f5fb8a097fe629f5a92fd528cf4...	HY484778		

5 rows × 30 columns



```
In [35]: 1 pd.set_option('display.max_rows', 100)
        2 df_people.isna().sum()
```

```
Out[35]: PERSON_ID          0
PERSON_TYPE          0
CRASH_RECORD_ID      0
RD_NO               10002
VEHICLE_ID           28251
CRASH_DATE           0
SEAT_NO             1145710
CITY                386612
STATE               372988
ZIPCODE             477431
SEX                 22243
AGE                417101
DRIVERS_LICENSE_STATE 592164
DRIVERS_LICENSE_CLASS 719487
SAFETY_EQUIPMENT      4087
AIRBAG_DEPLOYED       27110
EJECTION             17338
INJURY_CLASSIFICATION 632
HOSPITAL            1186950
EMS_AGENCY          1281208
EMS_RUN_NO          1411946
DRIVER_ACTION        294521
DRIVER_VISION        294925
PHYSICAL_CONDITION   293716
PEDPEDAL_ACTION      1410885
PEDPEDAL_VISIBILITY  1410940
PEDPEDAL_LOCATION    1410889
BAC_RESULT           293506
BAC_RESULT VALUE     1435858
CELL_PHONE_USE       1436454
dtype: int64
```

```
In [36]: 1 # Keeping Relevant Features
2 df_people_drop = df_people [[
3     'CRASH_RECORD_ID',
4     # 'RD_NO',
5     # 'PERSON_ID',
6     'PERSON_TYPE',
7     'VEHICLE_ID',
8     # 'CITY',
9     # 'STATE',
10    # 'ZIPCODE',
11    'SEX',
12    'AGE',
13    # 'DRIVERS_LICENSE_STATE',
14    # 'DRIVERS_LICENSE_CLASS',
15    'SAFETY_EQUIPMENT',
16    'AIRBAG_DEPLOYED',
17    # 'EJECTION',
18    # 'INJURY_CLASSIFICATION',
19    # 'DRIVER_VISION',
20    # 'DRIVER_ACTION',
21    # 'PHYSICAL_CONDITION',
22    # 'PEDPEDAL_ACTION',
23    # 'PEDPEDAL_VISIBILITY',
24    # 'PEDPEDAL_LOCATION',
25    # 'BAC_RESULT',
26    # 'BAC_RESULT VALUE',
27    # 'CELL_PHONE_USE',
28    ]]
29 print(df_people_drop.shape)
```

(1437611, 7)

```
In [37]: 1 # Filtering the data with Drivers only
2 df_people_driver = df_people_drop[df_people_drop.PERSON_TYPE == 'DRIVER']
3 print(df_people_driver.shape)
```

(1117947, 7)

```
In [38]: 1 df_people_driver.isna().sum()
```

```
Out[38]: CRASH_RECORD_ID      0
PERSON_TYPE      0
VEHICLE_ID      632
SEX      97
AGE      299295
SAFETY_EQUIPMENT      0
AIRBAG_DEPLOYED      0
dtype: int64
```

```
In [39]: 1 df_people_driver.describe()
```

```
Out[39]:
```

	VEHICLE_ID	AGE
count	1.117315e+06	818652.000000
mean	6.738542e+05	40.013475
std	3.924537e+05	15.836726
min	2.000000e+00	-177.000000
25%	3.338875e+05	27.000000
50%	6.705940e+05	37.000000
75%	1.014066e+06	51.000000
max	1.358762e+06	110.000000

```
In [40]: 1 # Dropping rows with driver age bigger than 90 and below 18
2 df_people_driver_age = df_people_driver[(df_people_driver['AGE'] >= 18) & (d
```

```
In [41]: 1 # Replacing all UNKNOWN values to missing values
2 #Unknown = ['UNKNOWN', 'USAGE UNKNOWN', 'DEPLOYMENT UNKNOWN']
3 df_vehicles_w_occ.replace({'Unknown' :np.nan , 'USAGE UNKNOWN' :np.nan , 'DE
```

<ipython-input-41-fad22d175994>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_vehicles_w_occ.replace({'Unknown' :np.nan , 'USAGE UNKNOWN' :np.nan , 'DE  
PLOYMENT UNKNOWN':np.nan }, inplace=True)
```

Merging Tables

```
In [42]: 1 merged = df_crashes_drop.merge(df_vehicles_w_occ, on='CRASH_RECORD_ID')
2 print(merged.shape)

(1160337, 18)
```

```
In [43]: 1 df = merged.merge(df_people_driver_age, on=['VEHICLE_ID', 'CRASH_RECORD_ID'])
2 print(df.shape)

(1099817, 23)
```

In [44]: 1 df.head()

Out[44]:

	CRASH_RECORD_ID	CRASH_DATE	POSTED_SPEED_LIMIT	WEA
0	062f5a6f6b87b762165d4da04d6d3a181385776a10b051...	08/31/2022 10:13:00 PM	30	
1	062f5a6f6b87b762165d4da04d6d3a181385776a10b051...	08/31/2022 10:13:00 PM	30	
2	0115ade9a755e835255508463f7e9c4a9a0b47e9304238...	07/15/2022 12:45:00 AM	30	
3	0115ade9a755e835255508463f7e9c4a9a0b47e9304238...	07/15/2022 12:45:00 AM	30	
4	017040c61958d2fa977c956b2bd2d6759ef7754496dc96...	07/15/2022 06:50:00 PM	30	

5 rows × 23 columns

In [45]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1099817 entries, 0 to 1099816
Data columns (total 23 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   CRASH_RECORD_ID             1099817 non-null object
1   CRASH_DATE                  1099817 non-null object
2   POSTED_SPEED_LIMIT          1099817 non-null int64
3   WEATHER_CONDITION           1055149 non-null object
4   ROADWAY_SURFACE_COND        1025920 non-null object
5   ROAD_DEFECT                 926974 non-null object
6   INJURIES_TOTAL              1099817 non-null float64
7   INJURIES_FATAL              1099817 non-null float64
8   CRASH_HOUR                  1099817 non-null int64
9   CRASH_DAY_OF_WEEK           1099817 non-null int64
10  CRASH_MONTH                  1099817 non-null int64
11  LATITUDE                    1092714 non-null float64
12  LONGITUDE                   1092714 non-null float64
13  has_injuries                 1099817 non-null int64
14  has_fatalities               1099817 non-null int64
```



```
In [46]: 1 # creating bins for times
2 # 0-6 = Late Night/Early Morning
3 # 6-12 = Morning
4 # 12-18 = Afternoon/Rush Hour
5 # 18-23 = Late Evening
6 df['time_bins'] = pd.cut(x=df['CRASH_HOUR'], bins = [-1,6,12,18,24],
7                        labels = ['Late Night/Early Morning',
8                                'Morning', 'Afternoon/Rush Hour', 'Late Evening'])
9 df.head(50)
```

```
Out[46]:
```

	CRASH_RECORD_ID	CRASH_DATE	POSTED_SPEED_LIMIT	V
0	062f5a6f6b87b762165d4da04d6d3a181385776a10b051...	08/31/2022 10:13:00 PM	30	
1	062f5a6f6b87b762165d4da04d6d3a181385776a10b051...	08/31/2022 10:13:00 PM	30	
2	0115ade9a755e835255508463f7e9c4a9a0b47e9304238...	07/15/2022 12:45:00 AM	30	
3	0115ade9a755e835255508463f7e9c4a9a0b47e9304238...	07/15/2022 12:45:00 AM	30	
4	017040c61958d2fa977c956b2bd2d6759ef7754496dc96...	07/15/2022 06:50:00 PM	30	
5	017040c61958d2fa977c956b2bd2d6759ef7754496dc96...	07/15/2022 06:50:00 PM	30	
6	01222759c6b8b0fd0f584226fbd88bdc549da3ed1e46255...	07/15/2022	40	

```
In [47]: 1 # Dropping Unnecassary Features
2 df_relv = df.drop([
3     'CRASH_RECORD_ID',
4     # 'RD_NO',
5     # 'PERSON_ID',
6     'CRASH_DATE',
7     'VEHICLE_ID',
8     # 'CITY',
9     # 'STATE',
10    # 'ZIPCODE',
11    'PERSON_TYPE',
12    'OCCUPANT_CNT',
13    # 'has_injuries',
14    # 'has_fatality',
15    # 'LONGITUDE',
16    # 'LATITUDE',
17    'ROAD_DEFECT',
18    # 'LIC_PLATE_STATE',
19    # 'TRAVEL_DIRECTION',
20    # 'DRIVERS_LICENSE_STATE',
21    # 'INJURY_CLASSIFICATION',
22    # 'DRIVER_ACTION',
23    # 'PHYSICAL_CONDITION'],
24    ],axis=1)
```

```
In [48]: 1 df_drop_missing = df_relv.dropna(subset = ['LATITUDE', 'LONGITUDE'], axis = 0
2         df_drop_missing.shape
```

Out[48]: (1092714, 18)

```
In [49]: 1 df_drop_missing.head()
```

Out[49]:

	POSTED_SPEED_LIMIT	WEATHER_CONDITION	ROADWAY_SURFACE_COND	INJURIES_TOTAL	IN
0	30	CLEAR	DRY	2.0	
1	30	CLEAR	DRY	2.0	
2	30	CLEAR	DRY	0.0	
3	30	CLEAR	DRY	0.0	
4	30	CLEAR	DRY	0.0	

< >

```
In [50]: 1 df1 = df_drop_missing[['LATITUDE', 'LONGITUDE']]
2         df1
```

Out[50]:

	LATITUDE	LONGITUDE
0	41.959389	-87.747348
1	41.959389	-87.747348
2	41.886336	-87.716203
3	41.886336	-87.716203
4	41.925111	-87.667997
...
1092709	41.735671	-87.663670
1092710	41.735671	-87.663670
1092711	41.877164	-87.720464
1092712	41.877164	-87.720464
1092713	41.989257	-87.776270

1092714 rows × 2 columns

```
In [51]: 1 # Binning Locations
2 from sklearn.cluster import KMeans
3 kmeans = KMeans(n_clusters = 50, init = 'k-means++')
4 kmeans.fit(df1[df1.columns]) # Compute k-means clustering.
5 df1['cluster_label'] = kmeans.fit_predict(df1[df1.columns])
6 #centers = kmeans.cluster_centers_ # Coordinates of cluster centers.
7 #Labels = kmeans.predict(df1[df1.columns]) # Labels of each point
8 df1.head(20)
```

<ipython-input-51-f4dd10255c2f>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df1['cluster_label'] = kmeans.fit_predict(df1[df1.columns])
```

```
Out[51]:
```

	LATITUDE	LONGITUDE	cluster_label
0	41.959389	-87.747348	15
1	41.959389	-87.747348	15
2	41.886336	-87.716203	39
3	41.886336	-87.716203	39
4	41.925111	-87.667997	43
5	41.925111	-87.667997	43
6	41.975826	-87.650420	9
7	41.975826	-87.650420	9
8	41.737337	-87.563560	19
9	41.737337	-87.563560	19
10	41.944199	-87.747157	31
11	41.944199	-87.747157	31
12	41.807712	-87.744440	7
13	41.855974	-87.663860	34
14	41.730216	-87.548387	41
15	41.956477	-87.785397	8
16	41.956477	-87.785397	8
17	41.807856	-87.733435	7
18	41.807856	-87.733435	7
19	41.877626	-87.629862	20

```
In [54]: 1 df_drop_missing['loc_clusters'] = df1['cluster_label']
```

```
In [55]: 1 df_drop_missing.isna().sum()
```

```
Out[55]: POSTED_SPEED_LIMIT      0
WEATHER_CONDITION      44507
ROADWAY_SURFACE_COND    73591
INJURIES_TOTAL          0
INJURIES_FATAL          0
CRASH_HOUR              0
CRASH_DAY_OF_WEEK       0
CRASH_MONTH             0
LATITUDE                0
LONGITUDE               0
has_injuries            0
has_fatality            0
VEHICLE_YEAR           199814
SEX                     97
AGE                    297434
SAFETY_EQUIPMENT        0
AIRBAG_DEPLOYED         0
time_bins               0
loc_clusters            0
dtype: int64
```

```
In [56]: 1 # drop missing rows
2 df_drop_missing2 = df_drop_missing.dropna(subset= ['AGE', 'VEHICLE_YEAR', 'ROA
3 print(df_drop_missing2.isna().sum())
4 print(df_drop_missing2.shape)
```

```
POSTED_SPEED_LIMIT      0
WEATHER_CONDITION      0
ROADWAY_SURFACE_COND    0
INJURIES_TOTAL          0
INJURIES_FATAL          0
CRASH_HOUR              0
CRASH_DAY_OF_WEEK       0
CRASH_MONTH             0
LATITUDE                0
LONGITUDE               0
has_injuries            0
has_fatality            0
VEHICLE_YEAR            0
SEX                     0
AGE                     0
SAFETY_EQUIPMENT        0
AIRBAG_DEPLOYED         0
time_bins               0
loc_clusters            0
dtype: int64
(715945, 19)
```

```
In [57]: 1 # Dropping outliers
2 counts = df_drop_missing2['SAFETY_EQUIPMENT'].value_counts()
3 df_drop_missing2 = df_drop_missing2[~df_drop_missing2['SAFETY_EQUIPMENT'].is
```

```
In [58]: 1 df_drop_missing2.SAFETY_EQUIPMENT.value_counts()
```

```
Out[58]: SAFETY BELT USED          467974
USAGE UNKNOWN          225687
NONE PRESENT           16523
SAFETY BELT NOT USED    3573
HELMET NOT USED         931
DOT COMPLIANT MOTORCYCLE HELMET    679
HELMET USED            391
NOT DOT COMPLIANT MOTORCYCLE HELMET 104
SHOULD/LAP BELT USED IMPROPERLY    79
Name: SAFETY_EQUIPMENT, dtype: int64
```

```
In [59]: 1 # Dropping outliers
2
3 counts = df_drop_missing2['WEATHER_CONDITION'].value_counts()
4 df_drop_missing2 = df_drop_missing2[~df_drop_missing2['WEATHER_CONDITION'].i
```

```
In [60]: 1 df_drop_missing2.to_csv('data/merged.csv')
```

```
In [419]: 1 # Sampling the population to make the data smaller
2 weights = {0 : 0.15, 1 : 0.85 }
3 df_drop_missing2['weights'] = df_drop_missing2['has_injuries'].apply(lambda
4 df_drop_missing2.head()
```

```
Out[419]:
```

	POSTED_SPEED_LIMIT	WEATHER_CONDITION	ROADWAY_SURFACE_COND	INJURIES_TOTAL	IN
--	--------------------	-------------------	----------------------	----------------	----

0	30	CLEAR	DRY	2.0
1	30	CLEAR	DRY	0.0
2	30	CLEAR	DRY	0.0
3	30	CLEAR	DRY	0.0
4	40	CLOUDY/OVERCAST	DRY	0.0

< >

```
In [420]: 1 # Sampling the population to make the data smaller
2
3 df_bal = df_drop_missing2.sample(frac = 0.1, weights='weights',random_state
4 df_bal.has_injuries.value_counts()
5
```

```
Out[420]: 0    36858
1     34736
Name: has_injuries, dtype: int64
```

```
In [421]: 1 for col in df_reduced.columns:
2         try:
3             print(col, df_reduced[col].value_counts(dropna=False)[:20])
4         except:
5             print(col, df_reduced[col].value_counts())
6             # If there aren't 5+ unique values for a column the first print stat
7             # will throw an error for an invalid idx slice
8             print('\n') # Break up the output between columns
```

```
POSTED_SPEED_LIMIT 30    54854
35    5462
25    3856
20    2079
15    1806
10    1060
40    869
0     631
45    536
5     316
55    46
3     23
50    19
39    8
9     6
65    3
2     3
34    3
60    3
~.    ~
```

Binning Imbalanced Features


Posted Speed Limit

```
In [422]: 1 # creating bins and label, previewing data
2 df_bal['speed_limit'] = pd.cut(x=df_bal['POSTED_SPEED_LIMIT'], bins = [-1,15
3                               labels = ['0-15', '16-25',
4                               '26-40', '41+'])
5 df_bal.speed_limit.value_counts()
```

```
Out[422]: 26-40    62641
16-25    5474
0-15    2820
41+    659
Name: speed_limit, dtype: int64
```

```
In [523]: 1 df_bal.groupby(['speed_limit']).INJURIES_TOTAL.sum().plot(kind='bar', color=
```

```
Out[523]: <AxesSubplot:xlabel='speed_limit'>
```




Vehicle Year

```
In [513]: 1 # creating bins and labels, preview data
2 df_bal['vehicle_year'] = pd.cut(x=df_bal['VEHICLE_YEAR'], bins = [0,2000,200
3                               labels = ['2000 & under', '2001-2005', '2006-2012', '
4
5 df_bal.vehicle_year.value_counts()
```

```
Out[513]: 2006-2012      22414
2015-2020      18033
2012-2015      14768
2001-2005      10815
2000 & under     4143
2021+          1245
Name: vehicle_year, dtype: int64
```

```
In [514]: 1 df_bal.groupby(['vehicle_year']).INJURIES_TOTAL.sum().plot(kind='bar', color
```

```
Out[514]: <AxesSubplot:xlabel='vehicle_year'>
```



AGE


```
In [499]: 1 # creating bins and labels, preview data
2 df_bal['age'] = pd.cut(x=df_bal['AGE'], bins = [0,25,35,55,100],
3 labels = ['24 & Under',
4 '25-35', '36-55', '56+'])
5 df_bal.age.value_counts()
6
```

```
Out[499]: 36-55          25191
25-35          19381
56+           13426
24 & Under     13420
Name: age, dtype: int64
```



```
In [500]: 1 df_bal.groupby(['age']).INJURIES_TOTAL.sum().plot(kind='bar', color=['#00CED1', '#FF69B4', '#FFD700', '#90EE90'])
```

```
Out[500]: <AxesSubplot: xlabel='age'>
```



Creating bins and labels, preview data

AIRBAG

In [425]:


```
1 # creating bins and labels, preview data
2 Airbag_map = {'DID NOT DEPLOY': 'DID NOT DEPLOY',
3               'NOT APPLICABLE': 'NOT APPLICABLE/UNKNOWN',
4               'DEPLOYMENT UNKNOWN': 'NOT APPLICABLE/UNKNOWN',
5               'DEPLOYED, FRONT': 'DEPLOYED',
6               'DEPLOYED, COMBINATION': 'DEPLOYED',
7               'DEPLOYED, SIDE': 'DEPLOYED',
8               'DEPLOYED OTHER (KNEE, AIR, BELT, ETC.)': 'DEPLOYED'}
9 df_bal.AIRBAG_DEPLOYED = df_bal.AIRBAG_DEPLOYED.map(Airbag_map)
10 print(df_bal.AIRBAG_DEPLOYED.value_counts())
11 #print(df_bal.Airbag.value_counts())
```

```
DID NOT DEPLOY          44373
NOT APPLICABLE/UNKNOWN  16305
DEPLOYED                10916
Name: AIRBAG_DEPLOYED, dtype: int64
```

In [520]:

```
1 df_bal.groupby(['AIRBAG_DEPLOYED']).INJURIES_TOTAL.sum().plot(kind='bar', co
```

Out[520]: <AxesSubplot:xlabel='AIRBAG_DEPLOYED'>



SAFETY EQ

```
In [426]: 1 safety_map = {'USAGE UNKNOWN': 'USAGE UNKNOWN',
2               'SAFETY BELT USED': 'SAFETY EQUIPMENT USED',
3               'NONE PRESENT': 'NONE PRESENT/USED',
4               'HELMET NOT USED': 'NONE PRESENT/USED',
5               'HELMET USED': 'SAFETY EQUIPMENT USED',
6               'SAFETY BELT NOT USED': 'NONE PRESENT/USED',
7               'NOT DOT COMPLIANT MOTORCYCLE HELMET ': 'NONE PRESENT/USED',
8               'DOT COMPLIANT MOTORCYCLE HELMET ': 'SAFETY EQUIPMENT USED',
9               'SHOULD/LAP BELT USED IMPROPERLY': 'NONE PRESENT/USED'}
10
11 df_bal.SAFETY_EQUIPMENT = df_bal.SAFETY_EQUIPMENT.map(safety_map)
12 df_bal.SAFETY_EQUIPMENT.value_counts()
```

```
Out[426]: SAFETY EQUIPMENT USED    44960
          USAGE UNKNOWN           24267
          NONE PRESENT/USED       2191
          Name: SAFETY_EQUIPMENT, dtype: int64
```

```
In [ ]: 1
```

Weather Condition

```
In [427]: 1 weather_map = {'CLEAR': 'CLEAR',
2               'RAIN': 'RAIN/CLOUDY/OTHER',
3               'CLOUDY/OVERCAST': 'RAIN/CLOUDY/OTHER',
4               'SNOW': 'RAIN/CLOUDY/OTHER',
5               'OTHER': 'RAIN/CLOUDY/OTHER',
6               'SLEET/HAIL': 'RAIN/CLOUDY/OTHER',
7               'FOG/SMOKE/HAZE': 'RAIN/CLOUDY/OTHER',
8               'FREEZING RAIN/DRIZZLE': 'RAIN/CLOUDY/OTHER',
9               'BLOWING SNOW': 'RAIN/CLOUDY/OTHER',
10              'SEVERE CROSS WIND GATE': 'RAIN/CLOUDY/OTHER'}
11
12 df_bal.WEATHER_CONDITION = df_bal.WEATHER_CONDITION.map(weather_map)
13 df_bal.WEATHER_CONDITION.value_counts()
```

```
Out[427]: CLEAR                58597
          RAIN/CLOUDY/OTHER    12997
          Name: WEATHER_CONDITION, dtype: int64
```

```
In [435]: 1 df_bal = df_bal.dropna( subset = 'SAFETY_EQUIPMENT', axis=0)
```

```
In [436]: 1 df_bal.isna().sum()
```

```
Out[436]: POSTED_SPEED_LIMIT      0
          WEATHER_CONDITION        0
          ROADWAY_SURFACE_COND     0
          INJURIES_TOTAL           0
          INJURIES_FATAL           0
          CRASH_HOUR               0
          CRASH_DAY_OF_WEEK        0
          CRASH_MONTH              0
          LATITUDE                 0
          LONGITUDE                0
          has_injuries              0
          has_fatality              0
          VEHICLE_YEAR             0
          SEX                      0
          AGE                      0
          SAFETY_EQUIPMENT          0
          AIRBAG_DEPLOYED          0
          time_bins                 0
          loc_clusters              0
          weights                   0
          speed_limit               0
          vehicle_year              0
          age                       0
          dtype: int64
```

```
In [ ]: 1
```


Visualization

```
In [827]: 1 plt.figure(figsize=(10,10))
          2 sns.countplot(x="CRASH_DAY_OF_WEEK", hue="has_injuries", data=df_bal)
```

...


```
In [852]: 1 plt.figure(figsize=(10,10))
          2 ax_time = sns.countplot(x="time_bins", hue="has_injuries", data=df_bal)
          3 ax_time.set(xlabel='TIME OF DAY', ylabel='COUNT')
```

```
Out[852]: [Text(0.5, 0, 'TIME OF DAY'), Text(0, 0.5, 'COUNT')]
```




```
In [840]: 1 plt.figure(figsize=(10,10))
          2 sns.countplot(x="age", hue="has_injuries", data=df_bal)
```

Out[840]: <AxesSubplot:xlabel='age', ylabel='count'>



```
In [846]: 1 plt.figure(figsize=(10,10))
          2 sns.countplot(y="WEATHER_CONDITION", hue="has_injuries", data=df_reduced)
```


Out[846]: <AxesSubplot:xlabel='count', ylabel='WEATHER_CONDITION'>



```
In [ ]: 1
```

```
In [851]: 1 plt.figure(figsize=(10,10))
          2 sns.countplot(x="AIRBAG_DEPLOYED", hue="has_injuries", data=df_bal)
```

```
Out[851]: <AxesSubplot:xlabel='AIRBAG_DEPLOYED', ylabel='count'>
```



```
In [ ]: 1 plt.figure(figsize=(10,10))
        2 sns.countplot(x="SEX", hue="has_injuries", data=df_bal)
```


In [700]:

```
1  # Splitting features into numeric and categorical
2  numeric_columns = [
3      'POSTED_SPEED_LIMIT',
4      # 'NUM_UNITS',
5      'VEHICLE_YEAR',
6      # 'OCCUPANT_CNT',
7      'AGE',
8
9  ]
10
11  cat_columns = [
12      'WEATHER_CONDITION',
13      'ROADWAY_SURFACE_COND',
14      # 'ROAD_DEFECT',
15      # 'CRASH_HOUR',
16      'CRASH_DAY_OF_WEEK',
17      'CRASH_MONTH',
18      # 'ZIPCODE',
19      'SEX',
20      'SAFETY_EQUIPMENT',
21      'AIRBAG_DEPLOYED',
22      'time_bins',
23      'speed_limit',
24      'age',
25      'vehicle_year',
26      # 'Airbag',
27
28  ]
29
30
31  Target1 = df_bal[[
32      'has_injuries',
33      # 'has_fatality',
34  ]]
35
36  Target2 = df_bal[[
37      # 'has_injuries',
38      'has_fatality',
39  ]]
40
41  numeric_df = df_bal[numeric_columns]
42  cat_df = df_bal[cat_columns]
```

Train_Test Split

```

In [1020]: 1 X = pd.concat([numeric_df,cat_df] , axis = 1)
           2 y = Target1
           3
           4 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = .25, r
           5
           6 X_train_nums = X_train [[
           7     #     'POSTED_SPEED_LIMIT',
           8     #     'VEHICLE_YEAR',
           9     #     'AGE',
          10
          11 ]]
          12
          13 X_train_cats = X_train [[
          14     'WEATHER_CONDITION',
          15     'ROADWAY_SURFACE_COND',
          16     #     'ROAD_DEFECT',
          17     #     'CRASH_HOUR',
          18     #     'CRASH_DAY_OF_WEEK',
          19     #     'CRASH_MONTH',
          20     #     'ZIPCODE',
          21     #     'SEX',
          22     #     'SAFETY_EQUIPMENT',
          23     #     'AIRBAG_DEPLOYED',
          24     #     'time_bins',
          25     'speed_limit',
          26     'age',
          27     #     'vehicle_year',
          28 ]]
          29
          30 X_train = pd.concat([X_train_nums,X_train_cats] , axis = 1)

```

```

In [449]: 1 # Cheking the distribution of targets
           2
           3 y_train.value_counts(normalize=True)

```

```

Out[449]: has_injuries
0          0.514889
1          0.485111
dtype: float64

```

Modelling

```

In [178]: 1 from imblearn.over_sampling import SMOTE
           2 from imblearn.under_sampling import RandomUnderSampler
           3 from imblearn.pipeline import Pipeline as imbPipeline
           4 from collections import Counter

```

```
In [179]: 1 counter = Counter(y_train)
2 print(counter)
3 print(y_train.value_counts(normalize=True))
```

```
Counter({'has_injuries': 1})
has_injuries
0          0.514778
1          0.485222
dtype: float64
```

```
In [1021]: 1 # One Hot Encoding categorical data
2
3
4 categorical_pipeline = Pipeline(steps=[
5     ('ohe', OneHotEncoder(drop='first',
6                           sparse=False))
7 ])
8
9 trans = ColumnTransformer(transformers=[
10     # ('numerical', numerical_pipeline, X_train_nums.columns),
11     ('categorical', categorical_pipeline, X_train_cats.columns)
12 ])
13
14 X_train_ohe = trans.fit_transform(X_train)
15 X_test_ohe = trans.transform(X_test)
```

C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\sklearn\compose_column_transformer.py:437: FutureWarning:

Given feature/column names or counts do not match the ones for the data given during fit. This will fail from v0.24.

Dummy Model

In [586]:

```
1 from sklearn.dummy import DummyClassifier
2 dummy_class = DummyClassifier(strategy = 'most_frequent')
3 dummy_class.fit(X_train,y_train)
4 #dummy_regr.predict(X_train)
5 y_dum = dummy_class.predict(X_test)
6 dummy_class.score(X_train, y_train)
7 dummy_class.score(X_test, y_test)
8 #dummy_train_RMSE = MSE(y_train,dummy_class.predict(X_train),squared = False)
9 #dummy_test_RMSE = MSE(y_test,dummy_class.predict(X_test), squared = False)
10
11 print(classification_report(y_test, y_dum))
```

	precision	recall	f1-score	support
0	0.52	1.00	0.68	9267
1	0.00	0.00	0.00	8588
accuracy			0.52	17855
macro avg	0.26	0.50	0.34	17855
weighted avg	0.27	0.52	0.35	17855

C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\sklearn\metrics_classification.py:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

First Model

```
In [1022]: 1 DT = DecisionTreeClassifier(random_state = 42, class_weight='balanced', max_
2 DT.fit(X_train_ohe, y_train)
3 y_DT = DT.predict(X_test_ohe)
4 classification_report(y_test, y_DT)
5
6 plot_confusion_matrix(DT, X_test_ohe, y_test)
7 for fi, feature in zip(DT.feature_importances_, X_train.columns):
8     print(fi, feature)
9 roc_auc_score(y_test, y_DT)
10
11 y_DT = DT.predict(X_test_ohe)
12
13 y_DT_proba = DT.predict_proba(X_test_ohe)[:, 1]
14 print(classification_report(y_test, y_DT))
15
16 print(DT.score(X_test_ohe, y_test))
17 print(DT.score(X_train_ohe, y_train))
```

```
0.03119341299074882 WEATHER_CONDITION
0.016854597790851588 ROADWAY_SURFACE_COND
0.012908849951574838 speed_limit
0.010801504264895604 age
      precision    recall  f1-score   support

      0         0.60      0.18      0.28       9267
      1         0.50      0.87      0.63       8588

 accuracy          0.51       17855
 macro avg         0.55      0.53      0.46       17855
 weighted avg      0.55      0.51      0.45       17855

0.5122934752170261
0.5241678024009111
```



Second Model

```
In [1023]: 1 LR_simple = LogisticRegression(max_iter = 1e3, random_state = 42, class_weig
2 # Create the GridSearchCV object with different hyperparameters
3
4
5 LR_simple.fit(X_train_ohe, y_train)
6
7
8 # Predict the label with the best model
9 y_pred_LR_simple = LR_simple.predict_proba(X_test_ohe)[: ,1]
10 y_pred_LR_simple_ = LR_simple.predict(X_test_ohe)
11
12 print(LR_simple.score(X_test_ohe, y_test))
13 print(LR_simple.score(X_train_ohe, y_train))
14
15 print(classification_report(y_test, y_pred_LR_simple_))
16 plot_confusion_matrix(LR_simple, X_test_ohe, y_test)
17 plot_roc_curve(LR_simple, X_test_ohe, y_test);
18 roc_auc_score(y_test, y_pred_LR_simple)
```

C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning:


A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().


0.51318958274993

0.5224501988312827

	precision	recall	f1-score	support
0	0.61	0.17	0.26	9267
1	0.50	0.89	0.64	8588
accuracy			0.51	17855
macro avg	0.55	0.53	0.45	17855
weighted avg	0.56	0.51	0.44	17855

Out[1023]: 0.5401740612011843





```
In [1015]: 1 # Logistic Regression Model with grid search
2
3 LR_pipe = LogisticRegression(max_iter = 1e3, random_state = 42, class_weight
4 # Create the GridSearchCV object with different hyperparameters
5 parameters_LR = {
6     'C': [1, 10, 100],
7     'solver': ['lbfgs', 'sag', 'saga']
8 }
9
10 cv_LR = GridSearchCV(LR_pipe, param_grid=parameters_LR, verbose = 2, n_jobs
11
12 cv_LR.fit(X_train_ohe, y_train)
13
14
15 # Predict the label with the best model
16 y_pred_LR = cv_LR.predict(X_test_ohe)
17 print(cv_LR.score(X_test_ohe, y_test))
18 print(cv_LR.score(X_train_ohe, y_train))
19 print(cv_LR.best_params_)
20 print(classification_report(y_test, y_pred_LR))
21 plot_confusion_matrix(cv_LR, X_test_ohe, y_test)
22 plot_roc_curve(cv_LR, X_test_ohe, y_test);
23 roc_auc_score(y_test, y_pred_LR)
```

Fitting 5 folds for each of 9 candidates, totalling 45 fits



```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:    0.6s
[Parallel(n_jobs=-1)]: Done  37 out of  45 | elapsed:    1.8s remaining:    0.3
s
[Parallel(n_jobs=-1)]: Done  45 out of  45 | elapsed:    2.1s finished
C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\skl
earn\utils\validation.py:72: DataConversionWarning:
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
0.51318958274993
0.5224501988312827
{'C': 1, 'solver': 'lbfgs'}
```

	precision	recall	f1-score	support
0	0.61	0.17	0.26	9267
1	0.50	0.89	0.64	8588
accuracy			0.51	17855
macro avg	0.55	0.53	0.45	17855
weighted avg	0.56	0.51	0.44	17855

Out[1015]: 0.526826438490994



```

In [1012]: 1 # DecisionTree Model with Pipeline
2 DT = DecisionTreeClassifier(random_state = 42, class_weight='balanced')
3
4 parameters_DT = {
5     'max_depth': [6,15,20],
6     'min_samples_split': [10,15],
7     'criterion': ['entropy', 'gini']
8 }
9
10 cv_DT = GridSearchCV(DT, param_grid=parameters_DT, verbose = 2, n_jobs = -1,
11
12 cv_DT.fit(X_train_ohe, y_train)
13
14 # Predict the label with the best model
15 y_pred_DT = cv_DT.predict(X_test_ohe)
16 print(cv_DT.score(X_test_ohe, y_test))
17 print(cv_DT.score(X_train_ohe, y_train))
18 print(cv_DT.best_params_)
19 print(classification_report(y_test, y_pred_DT))
20 plot_confusion_matrix(cv_DT, X_test_ohe, y_test)
21 plot_roc_curve(cv_DT, X_test_ohe, y_test);
22 roc_auc_score(y_test, y_pred_LR)

```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done   9 tasks      | elapsed:   0.0s
[Parallel(n_jobs=-1)]: Done  60 out of  60 | elapsed:   0.5s remaining:
0.0s
[Parallel(n_jobs=-1)]: Done  60 out of  60 | elapsed:   0.5s finished

```

0.5397563191433723

0.5509891530870029

```
{'criterion': 'gini', 'max_depth': 6, 'min_samples_split': 15}
```

	precision	recall	f1-score	support
0	0.60	0.17	0.27	9267
1	0.50	0.88	0.63	8588
accuracy			0.51	17855
macro avg	0.55	0.52	0.45	17855
weighted avg	0.55	0.51	0.44	17855

KNN Model

```
In [1011]: 1 # KNN Model with Pipeline
2 KNN = KNeighborsClassifier()
3
4 # Create the GridSearchCV object with different hyperparameters
5 parameters_KNN = {
6     'n_neighbors': [3, 5],
7     'metric': ['minkowski', 'manhattan'],
8     'weights': ['uniform'],
9 }
10
11 cv_KNN = GridSearchCV(KNN, param_grid=parameters_KNN, verbose = 2, n_jobs =
12
13 cv_KNN.fit(X_train_ohe, y_train)
14
15 # Predict the label with the best model
16 y_pred_KNN = cv_KNN.predict(X_test_ohe)
17 print(cv_KNN.score(X_test_ohe, y_test))
18 print(cv_KNN.score(X_train_ohe, y_train))
19 print(cv_KNN.best_params_)
20 print(classification_report(y_test, y_pred_KNN))
21 plot_confusion_matrix(cv_KNN, X_test_ohe, y_test)
22 plot_roc_curve(cv_KNN, X_test_ohe, y_test);
23 roc_auc_score(y_test, y_pred_LR)
```

Fitting 5 folds for each of 4 candidates, totalling 20 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done 11 out of 20 | elapsed: 33.6s remaining: 27.5
s
[Parallel(n_jobs=-1)]: Done 20 out of 20 | elapsed: 40.1s finished
C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\skl
earn\model_selection\_search.py:765: DataConversionWarning:
```



A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
0.5121193132936767
0.5209096479133745
{'metric': 'minkowski', 'n_neighbors': 5, 'weights': 'uniform'}
      precision    recall  f1-score   support

         0         0.53         0.60         0.56         9267
         1         0.49         0.42         0.45         8588

 accuracy          0.51          0.51          0.51         17855
 macro avg          0.51          0.51          0.51         17855
weighted avg          0.51          0.51          0.51         17855
```

Out[1011]: 0.526826438490994



Stacked Model

```
In [1010]: 1 LR_Stack = LogisticRegression(max_iter = 1e3, random_state = 42, class_weight='balanced')
2 DT_Stack = DecisionTreeClassifier(random_state = 42, class_weight='balanced')
3 KNN_Stack = KNeighborsClassifier(n_neighbors=5, metric='minkowski', weights='distance')
4
5 avg = VotingClassifier(estimators=[
6     ('LR', LR_Stack),
7     ('KNN', KNN_Stack),
8     ('DT', DT_Stack)])
9 avg.fit(X_train_ohe, y_train)
10
11 y_pred_avg = avg.predict(X_test_ohe)
12 print(avg.score(X_test_ohe, y_test))
13 print(avg.score(X_train_ohe, y_train))
14 print(classification_report(y_test, y_pred_avg))
15 plot_confusion_matrix(avg, X_test_ohe, y_test)
16 #plot_roc_curve(avg, X_test_ohe, y_test);
17 roc_auc_score(y_test, y_pred_avg)
18
```

C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning:


A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

0.5124614953794455

0.5237010623004686

	precision	recall	f1-score	support
0	0.60	0.19	0.28	9267
1	0.50	0.87	0.63	8588
accuracy			0.51	17855
macro avg	0.55	0.53	0.46	17855
weighted avg	0.55	0.51	0.45	17855

Out[1010]: 0.5253998253640674



XGBoost Model

```
In [1005]: 1 boost = xgboost.XGBClassifier(n_estimators = 1000, random_state=42, learning
2
3 boost.fit(X_train_ohe, y_train)
4
5 y_pred_boost = boost.predict(X_test_ohe)
6 print(boost.score(X_test_ohe, y_test))
7 print(boost.score(X_train_ohe, y_train))
8 print(classification_report(y_test, y_pred_boost))
9 plot_confusion_matrix(boost, X_test_ohe, y_test)
10 #plot_roc_curve(avg, X_test_ohe, y_test);
11 roc_auc_score(y_test, y_pred_boost)
12
```

C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\sklearn\utils\validation.py:72: DataConversionWarning:


A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

0.5269672360683282

0.5337266396579728

	precision	recall	f1-score	support
0	0.53	0.72	0.61	9267
1	0.51	0.31	0.39	8588
accuracy			0.53	17855
macro avg	0.52	0.52	0.50	17855
weighted avg	0.52	0.53	0.51	17855

Out[1005]: 0.5191538741800025



Random Forrest

```
In [1009]: 1 # Random Forrest Model with Pipeline
2 RF_pipe = imbPipeline(steps=[
3     ('RF', RandomForestClassifier(random_state = 42, class_weight='balanced'
4 ])
5
6 # Create the GridSearchCV object with different hyperparameters
7 parameters_RF = {
8     'RF__n_estimators': [100],
9     'RF__max_depth': [6, 8, 10, 20],
10    'RF__min_samples_split': [10, 15],
11    'RF__criterion': ['gini']
12 }
13
14 cv_RF = GridSearchCV(RF_pipe, param_grid=parameters_RF, verbose = 2, n_jobs
15
16 cv_RF.fit(X_train_ohe, y_train)
17
18 # Predict the label with the best model
19 y_pred_RF = cv_RF.predict_proba(X_test_ohe)
20 print(cv_RF.score(X_test_ohe, y_test))
21 print(cv_RF.score(X_train_ohe, y_train))
22 print(cv_RF.best_params_)
23 #print(classification_report(y_test, y_pred_RF))
24 plot_confusion_matrix(cv_RF, X_test_ohe, y_test)
25 #roc_auc_score(y_test, y_pred_RF)
```

Fitting 5 folds for each of 8 candidates, totalling 40 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 16 concurrent workers.
[Parallel(n_jobs=-1)]: Done   9 tasks   | elapsed:    2.3s
[Parallel(n_jobs=-1)]: Done  30 out of  40 | elapsed:    5.1s remaining:
1.6s
[Parallel(n_jobs=-1)]: Done  40 out of  40 | elapsed:    6.4s finished
C:\Users\milad\Documents\Flatiron\Anaconda\envs\learn-env\lib\site-packages\imblearn\pipeline.py:281: DataConversionWarning:
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
0.5355559042812542
0.550346051901595
{'RF__criterion': 'gini', 'RF__max_depth': 6, 'RF__min_samples_split': 15, 'R
F__n_estimators': 100}
```

Out[1009]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x26704ed80